

CS398-Deep Learning
Homework 1
Lingyi Xu (lingyix2)

1. Description of Implementation

First, I load the MNIST data from the dataset and preprocess the labels to one-hot format.

Fetch Data

```
1 # code from lecture notes
2 MNIST_data = h5py.File('MNISTdata.hdf5', 'r')
3 x_train = np.float32(MNIST_data['x_train'][:])
4 y_train = np.int32(np.array(MNIST_data['y_train'][:,0]))
5 x_test = np.float32(MNIST_data['x_test'][:])
6 y_test = np.int32(np.array(MNIST_data['y_test'][:,0]))
```

Helper Function for Data Pre-Process

```
1 def one_hot(Y):
2     m = Y.shape[0]
3     OHX = scipy.sparse.csr_matrix((np.ones(m), (Y, np.array(range(m)))))
4     OHX = np.array(OHX.todense()).T
5     return OHX

1 one_hot_y_test = one_hot(y_test)
2 one_hot_y_train = one_hot(y_train)
```

Then I build my mini-batch softmax logistic regression model.

I define the softmax function I will use to classify the probabilities:

$$F_{\text{softmax}}(z) = \frac{1}{\sum_{k=0}^{K-1} e^{z_k}} (e^{z_0}, e^{z_1}, \dots, e^{z_{K-1}}),$$

```
1 def softmax(z):
2     exp = np.exp(z - np.max(z, axis=1).reshape((-1,1)))
3     norms = np.sum(exp, axis=1).reshape((-1,1))
4     return exp / norms
```

And the gradient I will use for update W:

When training, I first initialize W, a zero matrix with shape (784,10). Then I use 2800 iterations to update the W by gradient decent.

For each iteration, I calculate the gradient of F(w) from the sampling:

$$\nabla_{\theta_j} J(\theta) = -\frac{1}{m} \sum_{i=1}^m [x^{(i)} (1\{y^{(i)} = j\} - p(y^{(i)} = j|x^{(i)}; \theta))]$$

```
1 def gradient(w,x,y):
2     n = [random.randint(0,len(x)-1) for i in range(batch_size)]
3     xn = x[n]
4     yn = y[n]
5     z = softmax(np.dot(xn,w))
6     gradient = -(1/batch_size)*np.dot(xn.T, (yn-z))
7     return gradient
```

After calculating the gradient of the function, we can update the parameter matrix w by (where “theta” is “W” in my implementation)

$$\theta^{(\ell+1)} = \theta^{(\ell)} - \alpha^{(\ell)} G^{(\ell)},$$

```
for i in range(total_iteration):
    gradient_i = gradient(w,x_train,y_train_cag)
    lr = learning_rate_arr[(i//int(total_iteration/4))]
    w -= (lr* gradient_i)
```

For this particular dataset, I use the following parameters:

Iteration = 2800

Batch size = 32

Learning rate = 0.5 when $0 \leq \text{itr} < 700$;
0.2 when $700 \leq \text{itr} < 1400$;
0.05 when $1400 \leq \text{itr} < 2100$;
0.01 when $2100 \leq \text{itr} < 2800$

```
1 batch_size = 32
2 total_iteration = 2800
3 learning_rate_arr = [0.5,0.2,0.05,0.01]
```

2. Final Test Accuracy:

91.9%

