

CS398-Deep Learning
Homework 2
Lingyi Xu (lingyix2)

1. Description of Implementation

First, I load the MNIST data from the dataset and preprocess the labels to one-hot format.

Fetch Data

```
1 # code from lecture notes
2 MNIST_data = h5py.File('MNISTdata.hdf5', 'r')
3 x_train = np.float32(MNIST_data['x_train'][:])
4 y_train = np.int32(np.array(MNIST_data['y_train'][:], 0))
5 x_test = np.float32(MNIST_data['x_test'][:])
6 y_test = np.int32(np.array(MNIST_data['y_test'][:], 0))
```

One_Hot Function for Data Pre-Process

```
1 def one_hot(Y):
2     m = Y.shape[0]
3     OHX = scipy.sparse.csr_matrix((np.ones(m), (Y, np.array(range(m)))
4     OHX = np.array(OHX.todense()).T
5     return OHX
6
7 one_hot_y_test = one_hot(y_test)
8 one_hot_y_train = one_hot(y_train)
```

Then I build my mini-batch softmax logistic regression model.
I define the nonlinearities $\sigma(z)$ (relu) function I will use:

```
1 def relu(x):
2     return np.maximum(0, x)
```

The forward propagation step:

Forward Propagation

```
1 def forward_prop(x, w1, w2):
2     hidden_output = relu(np.dot(x, w1))
3     output = relu(np.dot(hidden_output, w2))
4     return (hidden_output, output)
```

And the gradient I will use for update weight matrix w:

When training, I first initialize w1 and w2 to random value, w1 is the matrix which get x as input and hidden layer as output and w2 is the matrix which get hidden layer as input and y as output. Then I use 14000 iterations to update the W by gradient decent.

For each iteration, we can update the parameter matrix w by

$$W^{(\ell+1)} = W^{(\ell)} - \alpha^\ell \left(\delta \odot \sigma'(Z) \right) X^T$$

```
def train(self):
    #Randomly choose a batch from dataset
    n = [random.randint(0, len(self.X)-1) for i in range(self.batch_size)]
    batch_data = self.X[n]
    batch_label = self.y[n]
    #calculate the result of each layer
    hidden_output, output = forward_prop(batch_data, self.w1, self.w2)
    #calculate gradient of w2
    delta = output - batch_label
    grads_w2 = hidden_output.T.dot(delta)
    #calculate gradient of w1
    delta = (hidden_output > 0) * delta.dot(self.w2.T)
    grads_w1 = batch_data.T.dot(delta)
    #update w1 and w2
    self.w1 -= lr*(grads_w1/(self.batch_size))
    self.w2 -= lr*(grads_w2/(self.batch_size))
```

For this particular dataset, I use the following parameters:

Iteration = 14000

Batch size = 30


Learning rate = $0.25 / (0.001 * itr + 1)$

2. Final Test Accuracy:

96.5%

Plot the Accuracy vs. Iteration

```
1 plt.plot(accs)
2 plt.xlabel("iterations * 100")
3 plt.ylabel("accuracy")
4 plt.title("accuracy vs. iteration")
5 plt.show()
```



Print the Final Accuracy

```
1 print("final accuracy of test set is:", accs[-1])
```

final accuracy of test set is: 0.965