

CS398-Deep Learning

Homework 4

Lingyi Xu (lingyix2)

1. Description of Implementation

First, I implement the deep convolutional neural network following the architecture of lecture 6 slides.

Example of convolution network for CIFAR10:

- Convolution layer 1: 64 channels, $k = 4, s = 1, P = 2$.
- Batch normalization
- Convolution layer 2: 64 channels, $k = 4, s = 1, P = 2$.
- Max Pooling: $s = 2, k = 2$.
- Dropout
- Convolution layer 3: 64 channels, $k = 4, s = 1, P = 2$.
- Batch normalization
- Convolution layer 4: 64 channels, $k = 4, s = 1, P = 2$.
- Max Pooling \rightarrow Dropout

(Continued)

- Convolution layer 5: 64 channels, $k = 4, s = 1, P = 2$.
- Batch normalization
- Convolution layer 6: 64 channels, $k = 3, s = 1, P = 0$.
- Dropout
- Convolution layer 7: 64 channels, $k = 3, s = 1, P = 0$.
- Batch normalization
- Convolution layer 8: 64 channels, $k = 3, s = 1, P = 0$.
- Batch normalization, Dropout
- Fully connected layer 1: 500 units.
- Fully connected layer 2: 500 units.
- Linear \rightarrow Softmax function

```
def forward(self, x, extract_features=False):
    x = F.relu(self.bn1(self.conv1(x)))
    x = F.relu(self.conv2(x))
    x = F.max_pool2d(x, kernel_size=2, stride=2)
    x = self.drop1(x)
    x = F.relu(self.bn2(self.conv3(x)))
    x = F.relu(self.conv4(x))
    x = F.max_pool2d(x, kernel_size=2, stride=2)
    x = F.relu(self.bn3(self.conv5(x)))
    x = F.relu(self.conv6(x))
    x = self.drop3(x)
    x = F.relu(self.bn4(self.conv7(x)))
    x = F.relu(self.bn5(self.conv8(x)))
    x = self.drop4(x)
    x = x.view(-1, 4*4*64)
    x = F.relu(self.fc1(x))
    x = self.fc2(x)
    return x
```

Then I load and prepare the training data and test data of CIFAR10. When loading data, I used Transforms to normalize and enlarge the dataset to help training. For Transforms, I used random flip, random resize and random color jitter.

```
transform_train = transforms.Compose([
    transforms.RandomResizedCrop(DIM, scale=(0.7, 1.0), ratio=(1.0, 1.0)),
    transforms.ColorJitter(
        brightness = 0.1*torch.randn(1),
        contrast = 0.1*torch.randn(1),
        saturation = 0.1*torch.randn(1),
        hue = 0.1*torch.randn(1)),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

transform_test = transforms.Compose([
    transforms.CenterCrop(DIM),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

trainset = torchvision.datasets.CIFAR10(root='./', train=True, download=True, transform=transform_train)
testset = torchvision.datasets.CIFAR10(root='./', train=False, download=False, transform=transform_test)

trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle=True, num_workers=8)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=False, num_workers=8)
```

Train the model with GPU acceleration. For this particular dataset, I use the following parameters:

Iteration = 30, Batch size = 120, Learning rate = 0.001/ (0.1*itr+1)

```

for epoch in range(0,num_epochs):

    for param_group in optimizer.param_groups:
        param_group['lr'] = learning_rate/(1+epoch*0.1)

    train_accu = []
    model.train()

    for batch_idx, (X_train_batch,Y_train_batch) in enumerate(trainloader):

        if(Y_train_batch.shape[0]<batch_size):
            continue

        X_train_batch = Variable(X_train_batch).cuda()
        Y_train_batch = Variable(Y_train_batch).cuda()
        output = model(X_train_batch)

        loss = criterion(output, Y_train_batch)
        optimizer.zero_grad()

        loss.backward()
        optimizer.step()

        prediction = output.data.max(1)[1]
        accuracy = (float(prediction.eq(Y_train_batch.data).sum())/float(batch_size))*100.0

        train_accu.append(accuracy)
    accuracy_epoch = np.mean(train_accu)
    print(epoch,accuracy_epoch)

```

In the end, apply Monte Carlo methods, heuristic test methods, and normal test methods:

```

# Heuristic Prediction Rule
model.eval()
# Normal Prediction
with torch.no_grad():
    test_accu = []
    for batch_idx, (x_test_batch,y_test_batch) in enumerate(testloader):
        X_test_batch, Y_test_batch = Variable(x_test_batch).cuda(),Variable(y_test_batch).cuda()
        output = model(X_test_batch)
        prediction = output.data.max(1)[1]
        accuracy = (float(prediction.eq(Y_test_batch.data).sum())/float(batch_size))*100.0
        test_accu.append(accuracy)
    accuracy_test = np.mean(test_accu)
    print("Testing",accuracy_test)

#-----when using monte carlo prediction, use this part
# # Monte Carlo Prediction
# with torch.no_grad():
#     test_accu = []
#     for batch_idx, (X_test_batch,Y_test_batch) in enumerate(testloader):
#         X_test_batch, Y_test_batch = Variable(X_test_batch).cuda(),Variable(Y_test_batch).cuda()
#         prediction_vector = Variable(torch.tensor(np.zeros((50,120,10)))).cuda()
#         for i in range(monte_carlo_size):
#             output = model(X_test_batch)
#             temp_prediction = output.data
#             prediction_vector[i] = temp_prediction
#         prediction = (prediction_vector.mean(0)).max(1)[1]
#         accuracy = (float(prediction.eq(Y_test_batch.data).sum())/float(batch_size))*100.0
#         test_accu.append(accuracy)
#     mc_test_acc = np.mean(test_accu)
#     print("Monte Carlo test accuracy:",mc_test_acc)

```

2. Final Test Accuracy:

Normal: 83.9369%

Monte Carlo Method: 84.9326%

Heuristic Method: 84.0257%

3. Structure of CNN (from lecture slides):

Example of convolution network for CIFAR10:

- Convolution layer 1: 64 channels, $k = 4, s = 1, P = 2$.
- Batch normalization
- Convolution layer 2: 64 channels, $k = 4, s = 1, P = 2$.
- Max Pooling: $s = 2, k = 2$.
- Dropout
- Convolution layer 3: 64 channels, $k = 4, s = 1, P = 2$.
- Batch normalization
- Convolution layer 4: 64 channels, $k = 4, s = 1, P = 2$.
- Max Pooling → Dropout

(Continued)

- Convolution layer 5: 64 channels, $k = 4, s = 1, P = 2$.
- Batch normalization
- Convolution layer 6: 64 channels, $k = 3, s = 1, P = 0$.
- Dropout
- Convolution layer 7: 64 channels, $k = 3, s = 1, P = 0$.
- Batch normalization
- Convolution layer 8: 64 channels, $k = 3, s = 1, P = 0$.
- Batch normalization, Dropout
- Fully connected layer 1: 500 units.
- Fully connected layer 2: 500 units.
- Linear → Softmax function